

Final Project Report

Prepared by,

1. G.Araf Ahmed (22101532) 2.

Md Shihab Sarker (22101516) 3.

Fahim Foysal Abit (22101543)

Course Code: CSE440

Section: 3

Group No: 1

Final Project Report

Abstract:

This project runs different Machine Learning and Neural Network models with different vectorization methods and does experiments to find the best models and techniques. First we loaded the datasets in the code environment then did extensive exploratory data analysis. The EDA shows us there are a total 280K rows and 2 columns in both the test and train dataset. There are 10 categories in the target feature “Class”. The “QA Text” column had three inputs in it, question title, question content and best answer. We found that merging all these types into a plain text give the best result so we merged them into a single text string and did data preprocessing on them. Then we did data preprocessing using lemmatization and then split the train data set in train and validation in a 80-20 ratio. Then we vectorized the data using different vectorization methods, TF-IDF, BOW, Glove and Skip gram. And we ran different machine learning and neural network models on each of them totalling 22 separate experiments. We found out that the best model for BOW vectorization is Naive Bayes, for TF-IDF Logistic Regression, for both Glove and Skip gram GRU.

Introduction:

The task was given as a final project for the NLP course CSE440. We were given 2 datasets. One being the test dataset and the other train. Both of them had a total

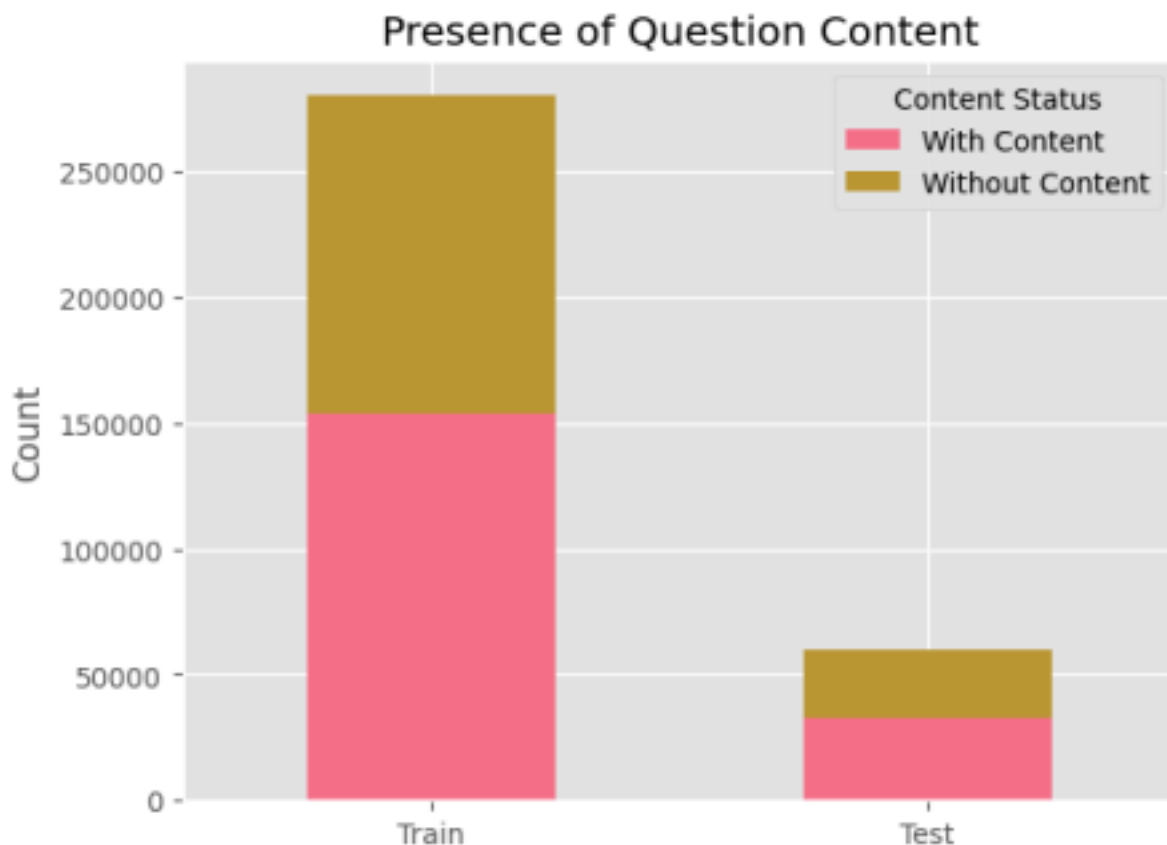
280k rows and 2 columns each. The target feature “Class” had 10 unique categories. They were almost evenly distributed throughout the dataset. The “QA Text” column had three components. Question Title, Question Content and Best Answer. We found that the average text length pattern for title length is 10.7 words, average answer length: 59.2 words and content is often missing: 45.0% of cases.

Methodology:

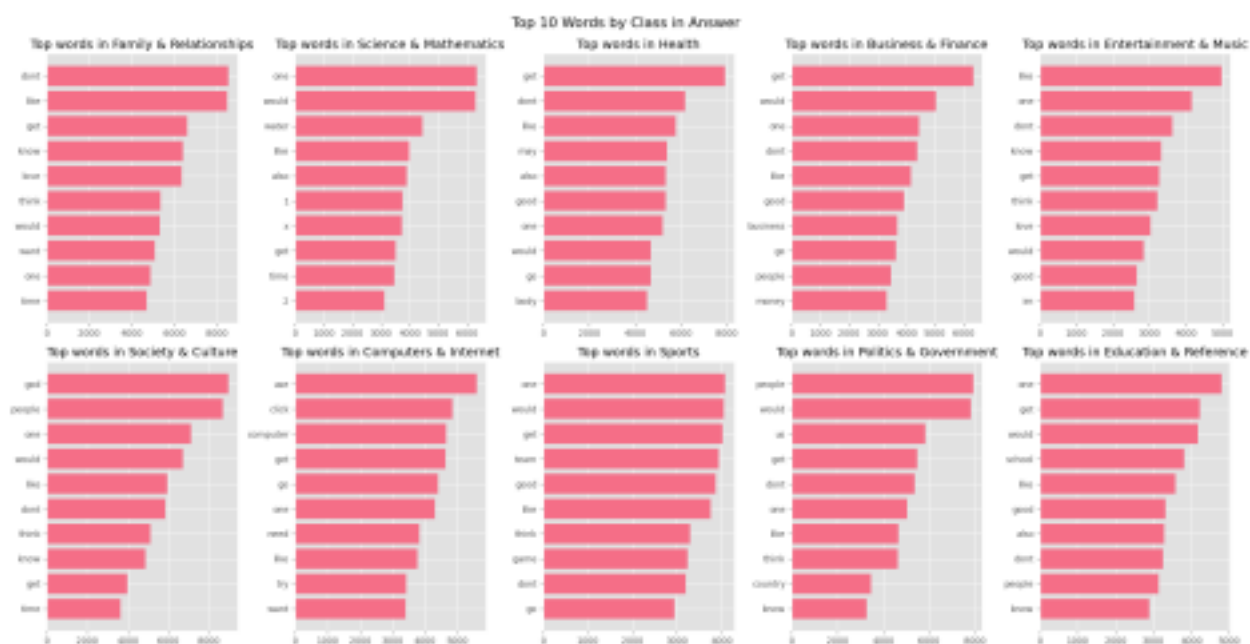
While doing exploratory data analysis we found that we had 10 unique categories in our target feature “Class”. We then ran code to see the class distributions. The class distribution for our both datasets among these 10 categories were almost similar. Below is the histogram for both of the Class columns from train dataset and test dataset.



We found that the Question Content component in the “QA Test” column has a lot of missing data. We just replaced the Nan values with empty strings.



We also ran codes to see the top 10 words in Question title, Question Component, Best Answer and the target feature “Class”. Below is the graph for top 10 words in the Class feature.

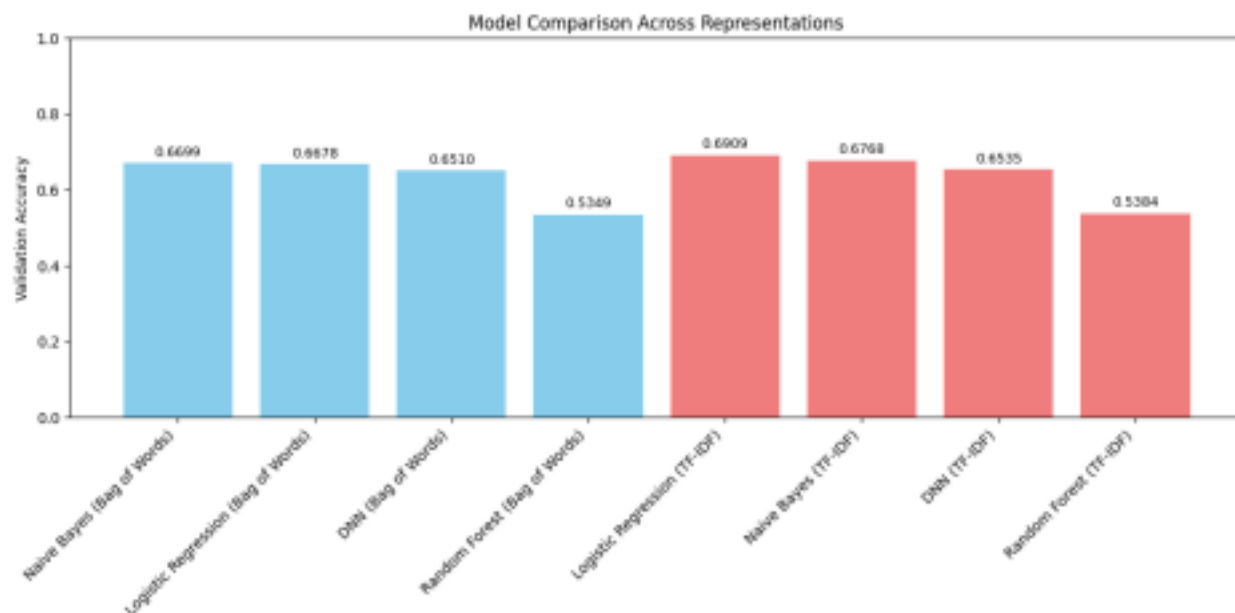


We ran the data on a Logistic regression model using TF-IDF vectorization to extract information on how we should create the vectors. We found out concatenating all those components into a single text string gives better output than separately vectorizing those components and using hstack on them later. That's why we concatenated the components. Then we removed all the punctuation marks, turned the text into lower case handled Nan values with empty strings and then removed stopwords. Then we ran lemmatization techniques on the data. Now we have clean data ready for vectorization. First we used 2 vectorization methods, BOW and TF-IDF and after vectorizing the data for each of them we ran the vectors on the Logistic Regression model, Naive Bayes model, Random forest model and Deep Neural Network model. We used different parameter sets to see which parameters gives the best results. With this we were done doing $2 \times 4 = 8$ experiments. Next we used Glove vectorization with 7 Neural Network models. We used Deep Neural Network, SimpleRNN, GRU, LSTM, Bidirectional SimpleRNN, Bidirectional GRU and Bidirectional LSTM. Then we used Skip Gram vectorization to run the same 7 models. We used a 3 hidden layer deep neural network for both BOW and TF-IDF vectorization. For Glove vectorization we used different neural network models with different hyper parameters. For the Skipgram vectorization we also use all those 7 Neural Network models. The Deep Neural Network model had only one hidden layers due to it taking too much time to run if we increased the layers. We used 512 batch size for it and 15 epochs. And for other models we used different batch sizes and epochs based mainly on saving while running the models smoothly.

Results:

For the BOW vectorization among Logistic Regression, Naive Bayes, Random Forest and Deep Neural Network, Naive Bayes did the best with a F1-macro score of 0.6622.

In the TF-IDF vectorization part among the same models Logistic Regression produced the best result. It had a F1-macro score of 0.6878.



And here is a screenshot of the performance of other models on their best parameter set:

```

=====
COMPREHENSIVE COMPARISON OF ALL MODELS ACROSS BOTH REPRESENTATIONS
=====
Detailed Results Table:
Representation      Model      Validation Accuracy  F1-weighted  F1-macro
Bag of words        Naive Bayes  0.669911          0.666162    0.665950
Bag of words        Logistic Regression  0.667768          0.643868    0.642090
Bag of words        DNN         0.651888          0.646598    0.646281
Bag of words        Random Forest  0.534875          0.534394    0.533340
TF-IDF              Logistic Regression  0.690891          0.688128    0.687789
TF-IDF              Naive Bayes  0.675823          0.672215    0.671903
TF-IDF              DNN         0.653500          0.653496    0.653254
TF-IDF              Random Forest  0.538375          0.536921    0.536470

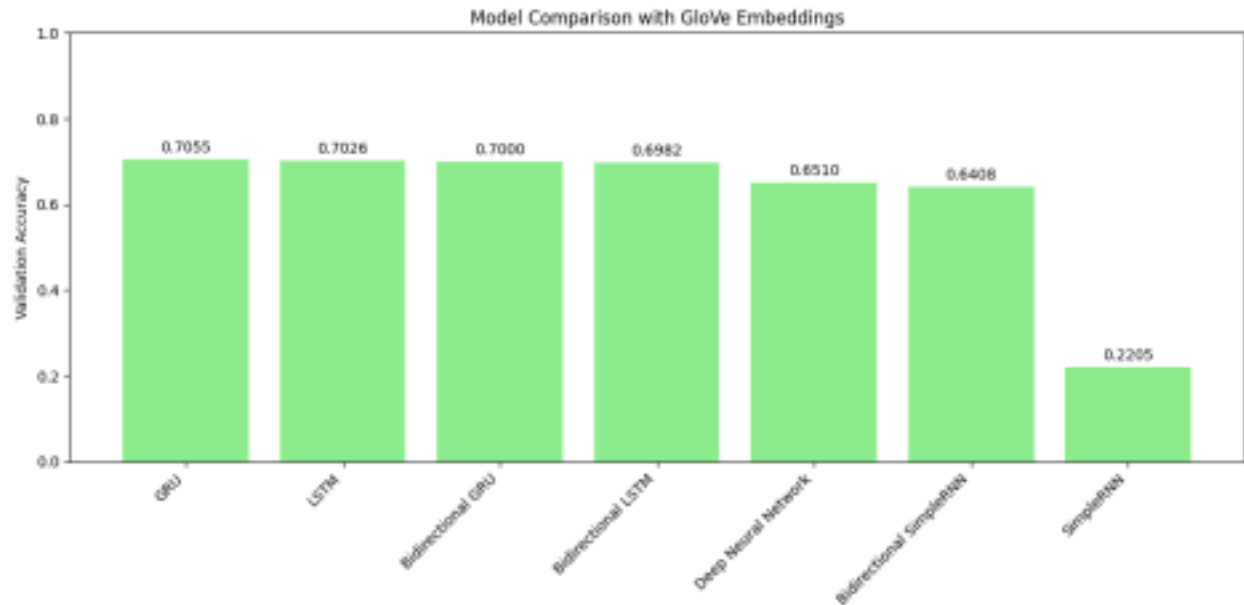
Parameters
[ 'alpha': 1.0]
[ 'C': 0.1, 'class_weight': 'balanced', 'max_iter': 500, 'n_jobs': -1, 'random_state': 440]
[ 'architecture': '512-256-128-64', 'dropout': '0.5-0.4-0.3-0.2']
[ 'max_depth': 20, 'n_estimators': 50, 'n_jobs': -1, 'random_state': 440]
[ 'C': 1.0, 'class_weight': 'balanced', 'max_iter': 500, 'n_jobs': -1, 'random_state': 440]
[ 'alpha': 0.5]
[ 'architecture': '512-256-128-64', 'dropout': '0.5-0.4-0.3-0.2']
[ 'max_depth': 20, 'n_estimators': 50, 'n_jobs': -1, 'random_state': 440]

```

For Glove vectorization among the 7 Neural Network model GRU produced the best score, an F1-macro of ~0.70. Below is the confusion matrix of GRU ran on Glove:



And the performance of other models on Glove is also below:



```
=====
COMPREHENSIVE COMPARISON OF ALL MODELS WITH GLOVE
EVALUATED ON VALIDATION SET
=====
```

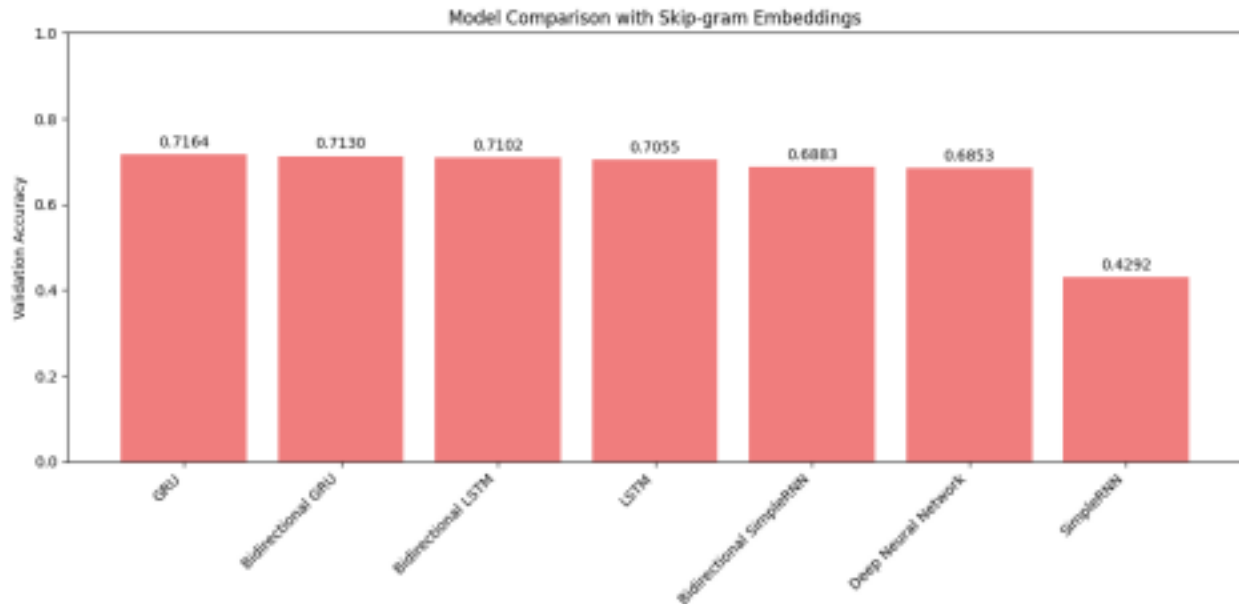
Detailed Results Table (Validation Set Performance):

Model	Validation Accuracy	F1-weighted	F1-macro	Training Time (s)
GRU	0.705482	0.699806	0.699468	164.833907
LSTM	0.702589	0.697659	0.697318	191.729682
Bidirectional GRU	0.699964	0.694995	0.694651	154.962499
Bidirectional LSTM	0.698196	0.693772	0.693431	164.724034
Deep Neural Network	0.650982	0.644148	0.643812	28.211925
Bidirectional SimpleRNN	0.640839	0.637093	0.636745	139.007995
SimpleRNN	0.220500	0.165516	0.165559	37.633789

For the Skipgram model GRU again produced the best result. An F1-macro of 0.71. Below is the confusion matrix GRU ran on embedding.



The performance of other models that ran on GRU is also below:



```
=====
COMPREHENSIVE COMPARISON OF ALL MODELS WITH SKIP-GRAM
EVALUATED ON VALIDATION SET
=====
```

Detailed Results Table (Validation Set Performance):

Model	Validation Accuracy	F1-weighted	F1-macro	Training Time (s)
GRU	0.716375	0.711213	0.710872	328.119132
Bidirectional GRU	0.713036	0.707583	0.707239	443.443572
Bidirectional LSTM	0.710196	0.706341	0.706014	438.005053
LSTM	0.705464	0.699570	0.699201	334.896129
Bidirectional SimpleRNN	0.688286	0.685018	0.684692	3850.674301
Deep Neural Network	0.685339	0.676339	0.675977	109.772571
SimpleRNN	0.429214	0.409913	0.409289	1046.960927

Among all the 22 experiments we ran we got the best performance from GRU that ran on Skipgram embedding. And the Simple RNN model ran on Glove produced the worst result with an F1-macro score of 0.165.

Conclusion:

To sum up, for BOW we had the best performance from Naive Bayes, for TF-IDF Logistic Regression, for both Glove and Skipgram we had GRU. GRU is clearly the winner here with the highest validation, F1-weighted and F1-macro scores. The Neural Network models on average produced the best results. Only the simple RNN model produced very poor results. The simple RNN model produced

the poorest result overall. Our main limitation in this project was time and resources. We had very little time and the models were taking a lot of time to train. And the low GPU usage time made things harder for us. That's why we had to make the models very simple, with few hidden layers and a big batch size. In the future we could produce better results if we have enough time and enough resources. We could make the Neural Network models a bit more complex, have a bit more epochs and achieve better results.

References

[1] J. Pennington, R. Socher, and C. D. Manning, "GloVe: Global vectors for word representation,"

EMNLP, 2014.

[2] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector

space," ICLR, 2013.

[3] Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR

2011. [4] F. Chollet et al., "Keras," 2015. [Online]. Available:

<https://keras.io>

[5] TensorFlow, Google Brain Team. [Online]. Available:

<https://www.tensorflow.org>

[6] NLTK: Natural Language Toolkit. [Online]. Available: <https://www.nltk.org>